

Title: Analyzing video games sales ¶

Information about the data used:

-Data Source: Video Game Sales Dataset

-Location: <https://www.kaggle.com/gregorut/videogamesales>

(<https://www.kaggle.com/gregorut/videogamesales>)

-Filename: videogamesales.zip

Initial the exploration of the Dataset

```
In [5]: import numpy as np
import pandas as pd
```

Importing data

```
In [6]: df = pd.read_csv('vgsales.csv')
```

```
In [7]: df.head()
```

```
Out[7]:
```

	Rank	Name	Platform	Year	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other
0	1	Wii Sports	Wii	2006.0	Sports	Nintendo	41.49	29.02	3.77	
1	2	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	3.58	6.81	
2	3	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.85	12.88	3.79	
3	4	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.75	11.01	3.28	
4	5	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	8.89	10.22	

```
In [9]: df.shape
```

```
Out[9]: (16598, 11)
```

```
In [10]: df.columns
```

```
Out[10]: Index(['Rank', 'Name', 'Platform', 'Year', 'Genre', 'Publisher', 'NA_Sales', 'EU_Sales', 'JP_Sales', 'Other_Sales', 'Global_Sales'], dtype='object')
```

```
In [12]: df.describe()
```

```
Out[12]:
```

	Rank	Year	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Globa
count	16598.000000	16327.000000	16598.000000	16598.000000	16598.000000	16598.000000	16598
mean	8300.605254	2006.406443	0.264667	0.146652	0.077782	0.048063	0
std	4791.853933	5.828981	0.816683	0.505351	0.309291	0.188588	1
min	1.000000	1980.000000	0.000000	0.000000	0.000000	0.000000	0
25%	4151.250000	2003.000000	0.000000	0.000000	0.000000	0.000000	0
50%	8300.500000	2007.000000	0.080000	0.020000	0.000000	0.010000	0
75%	12449.750000	2010.000000	0.240000	0.110000	0.040000	0.040000	0
max	16600.000000	2020.000000	41.490000	29.020000	10.220000	10.570000	82

Know we will clean our data but let us first check if there is any null value

```
In [39]: df.isnull().values.any()
```

```
Out[39]: True
```

```
In [40]: clean_df= df.dropna()
```

```
In [15]: clean_df.isnull().values.any()
```

```
Out[15]: False
```

Dimentions of new cleaned data

```
In [18]: clean_df.shape
```

```
Out[18]: (16291, 11)
```

```
In [19]: clean_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16291 entries, 0 to 16597
Data columns (total 11 columns):
Rank                16291 non-null int64
Name                16291 non-null object
Platform            16291 non-null object
Year                16291 non-null float64
Genre                16291 non-null object
Publisher            16291 non-null object
NA_Sales             16291 non-null float64
EU_Sales             16291 non-null float64
JP_Sales             16291 non-null float64
Other_Sales          16291 non-null float64
Global_Sales         16291 non-null float64
dtypes: float64(6), int64(1), object(4)
memory usage: 1.5+ MB
```


Importing seaborn and matplotlib

```
In [22]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Data Visulation


```
In [43]: plt.figure(figsize=(10,5))
sns.countplot(clean_df['Genre'])
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x24b4a8ba5c0>
```



```
In [44]: plt.figure(figsize=(15,5))
sns.countplot(df['Platform'])
```

Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x24b49d09400>




```
In [45]: platGenre = pd.crosstab(df.Platform,df.Genre)
platGenre.head(5)
```

Out[45]:

Genre	Action	Adventure	Fighting	Misc	Platform	Puzzle	Racing	Role-Playing	Shooter	Simulatio
Platform										
2600	61	2	2	5	9	11	6	0	24	
3DO	0	1	0	0	0	1	0	0	0	
3DS	182	37	14	53	28	20	11	86	7	3
DC	3	11	12	0	2	0	6	4	3	
DS	343	240	36	393	92	238	67	200	42	28

```
In [46]: platGenreTotal = platGenre.sum(axis=1).sort_values(ascending = False)
plt.figure(figsize=(10,15))
sns.barplot(x=platGenreTotal.values,y=platGenreTotal.index)
```


Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x24b49722cc0>



So You can see DS and PS2 have the most games in their platform and now, we can see the detail of genre of game in platform that have more than 1000 games. In the seaborn, it is easy to use heatmap rather than the stacked bar, so we can use the heatmap to have a try.


```
In [47]: pub=clean_df.groupby('Publisher')['Publisher'].count().sort_values(ascending = False)
sns.barplot(x=pub.values,y=pub.index)
```

Out[47]: <matplotlib.axes._subplots.AxesSubplot at 0x24b461e8278>




```
In [49]: yr=clean_df.groupby('Year')['Year'].count().sort_values(ascending = False).head(10)
plt.plot(yr)
```

Out[49]: [<matplotlib.lines.Line2D at 0x24b49c657b8>]



```
In [50]: plt.figure(figsize=(10,5))  
sns.pointplot(x=yr.index ,y=yr.values)
```

```
Out[50]: <matplotlib.axes._subplots.AxesSubplot at 0x24b49c43f60>
```



So it is showing most important year is 2009 based on others year

convert Categorical variable to numerical Variables

```
In [73]: from sklearn import model_selection, preprocessing
for c in clean_df.columns:
    if clean_df[c].dtype == 'object':
        lbl = preprocessing.LabelEncoder()
        lbl.fit(list(clean_df[c].values))
        df[c] = lbl.transform(list(clean_df[c].values))
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-73-73b8d0b09c90> in <module>
      4         lbl = preprocessing.LabelEncoder()
      5         lbl.fit(list(clean_df[c].values))
----> 6         df[c] = lbl.transform(list(clean_df[c].values))
      7

~\Anaconda3\lib\site-packages\pandas\core\frame.py in __setitem__(self, key, value)
    3117         else:
    3118             # set column
-> 3119             self._set_item(key, value)
    3120
    3121     def _setitem_slice(self, key, value):

~\Anaconda3\lib\site-packages\pandas\core\frame.py in _set_item(self, key, value)
    3192
    3193         self._ensure_valid_index(value)
-> 3194         value = self._sanitize_column(key, value)
    3195         NDFrame._set_item(self, key, value)
    3196

~\Anaconda3\lib\site-packages\pandas\core\frame.py in _sanitize_column(self, key, value, broadcast)
    3389
    3390         # turn me into an ndarray
-> 3391         value = _sanitize_index(value, self.index, copy=False)
    3392         if not isinstance(value, (np.ndarray, Index)):
    3393             if isinstance(value, list) and len(value) > 0:

~\Anaconda3\lib\site-packages\pandas\core\series.py in _sanitize_index(data, index, copy)
    3999
    4000         if len(data) != len(index):
-> 4001             raise ValueError('Length of values does not match length of '
    4002                                'index')
    4002
    4003         if isinstance(data, ABCIndexClass) and not copy:
```

ValueError: Length of values does not match length of index


```
In [74]: corr=clean_df.corr()
corr = (corr)
plt.figure(figsize=(14,14))
sns.heatmap(corr, cbar = True, square = True, annot=True, fmt= '.2f',annot_kws=
{'size': 15},
            xticklabels=corr.columns.values,
            yticklabels=corr.columns.values)
sns.plt.title('Heatmap of Correlation Matrix')
```

AttributeError Traceback (most recent call last)

```
<ipython-input-74-8c8a82c585d2> in <module>
      7             xticklabels=corr.columns.values,
      8             yticklabels=corr.columns.values)
----> 9 sns.plt.title('Heatmap of Correlation Matrix')
```

AttributeError: module 'seaborn' has no attribute 'plt'




so all sales are highly correlated with global sales which is very true..Rank is not showing highly correlated with global sales as this is shown in Matrix formation.

```
In [53]: plt.figure(figsize=(10,5))
sns.regplot(x='Rank',y='Global_Sales',data=clean_df)
```

C:\Users\Aymen\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.


```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[53]: <matplotlib.axes._subplots.AxesSubplot at 0x24b4cc18978>
```



```
In [55]: plt.figure(figsize=(10,5))
sns.regplot(x='NA_Sales',y='Global_Sales',data=clean_df)
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x24b4d507278>
```



```
In [56]: clean_df.dtypes
```

```
Out[56]: Rank          int64
Name          object
Platform      object
Year          float64
Genre         object
Publisher     object
NA_Sales      float64
EU_Sales      float64
JP_Sales      float64
Other_Sales   float64
Global_Sales   float64
dtype: object
```

```
In [57]: #Train-Test split
from sklearn.model_selection import train_test_split
label = clean_df.pop('Global_Sales')
data_train, data_test, label_train, label_test = train_test_split(clean_df, label,
```

```
In [58]: data_train.shape,data_test.shape
```

```
Out[58]: ((13032, 10), (3259, 10))
```

Start with Extream gradient boosting

```
In [65]: #!/pip install xgboost
```

Collecting xgboost

Downloading https://files.pythonhosted.org/packages/51/c1/198915b13e98b62a98f48309c41012638464651da755d941f4abe384c012/xgboost-0.82-py2.py3-none-win_amd64.whl (https://files.pythonhosted.org/packages/51/c1/198915b13e98b62a98f48309c41012638464651da755d941f4abe384c012/xgboost-0.82-py2.py3-none-win_amd64.whl) (7.7M B)

Requirement already satisfied: scipy in c:\users\aymen\anaconda3\lib\site-packages (from xgboost) (1.1.0)

Requirement already satisfied: numpy in c:\users\aymen\anaconda3\lib\site-packages (from xgboost) (1.15.4)

Installing collected packages: xgboost

Successfully installed xgboost-0.82

```
In [66]: import xgboost as xgb
```

```
from sklearn.model_selection import KFold, train_test_split, GridSearchCV
```

```
In [67]: xgb_params = {  
    'eta': 0.05,  
    'max_depth': 5,  
    'subsample': 0.7,  
    'colsample_bytree': 0.7,  
    'objective': 'reg:linear',  
    'eval_metric': 'rmse',  
    'silent': 1  
}
```

```
In [69]: dtrain = xgb.DMatrix(data_train, label_train)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-69-f2c2fb056c93> in <module>
----> 1 dtrain = xgb.DMatrix(data_train, label_train)

~\Anaconda3\lib\site-packages\xgboost\core.py in __init__(self, data, label, mi
ssing, weight, silent, feature_names, feature_types, nthread)
    382         data, feature_names, feature_types = _maybe_pandas_data(data,
    383                                                                    feature
_names,
--> 384                                                                    feature
_types)
    385
    386         data, feature_names, feature_types = _maybe_dt_data(data,

~\Anaconda3\lib\site-packages\xgboost\core.py in _maybe_pandas_data(data, featu
re_names, feature_types)
    239         msg = """DataFrame.dtypes for data must be int, float or bool.
    240                Did not expect the data types in fields """
--> 241         raise ValueError(msg + ', '.join(bad_fields))
    242
    243         if feature_names is None:

ValueError: DataFrame.dtypes for data must be int, float or bool.
                Did not expect the data types in fields Name, Platform, Genre,
                Publisher
```

```
In [70]: lr_data_train=data_train[['NA_Sales','EU_Sales','JP_Sales','Other_Sales']]
lr_data_test=data_test[['NA_Sales','EU_Sales','JP_Sales','Other_Sales']]
lr_label_train=label_train
lr_label_test=label_test
```

```
In [71]: #Linear Regression
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(lr_data_train, lr_label_train)
lr_score_train = lr.score(lr_data_train, lr_label_train)
print("Training score: ",lr_score_train)
lr_score_test = lr.score(lr_data_test, lr_label_test)
print("Testing score: ",lr_score_test)
```

```
Training score:  0.9999889271662358
Testing score:  0.9999887509682348
```

The accuracy improved to 99%

```
In [75]: y_pre = lr.predict(lr_data_test)
```

```
In [76]: out_lr = pd.DataFrame({'Actual_Global_Sales': lr_label_test, 'Predict_Global_Sales': lr_prediction_test})
out_lr[['Actual_Global_Sales', 'Predict_Global_Sales', 'Diff']].head(5)
```

Out[76]:

	Actual_Global_Sales	Predict_Global_Sales	Diff
9317	0.14	0.130333	0.009667
14835	0.03	0.020339	0.009661
9752	0.12	0.120335	-0.000335
10251	0.11	0.110324	-0.000324
16565	0.01	0.010337	-0.000337

```
In [77]: out_lr.shape
```


Out[77]: (3259, 3)

```
In [78]: sns.regplot(out_lr['Predict_Global_Sales'], out_lr['Diff'])
```

C:\Users\Aymen\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

Out[78]: <matplotlib.axes._subplots.AxesSubplot at 0x24b4cac4908>



```
In [79]: #Ensemble XGBOOST & LINEAR REGRESSOR for train data
en_dtest=xgb.DMatrix(data_train)
y_xgb_pred = model.predict(en_dtest)

y_lr_pred = lr.predict(lr_data_train)

Ensemble=pd.DataFrame({'XGBOOST':y_xgb_pred , 'LINEAR_REG':y_lr_pred , 'GLOBAL_SALI
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-79-d8edccb37eeb> in <module>
      1 #Ensemble XGBOOST & LINEAR REGRESSOR for train data
----> 2 en_dtest=xgb.DMatrix(data_train)
      3 y_xgb_pred = model.predict(en_dtest)
      4
      5 y_lr_pred = lr.predict(lr_data_train)

~\Anaconda3\lib\site-packages\xgboost\core.py in __init__(self, data, label, missing, weight, silent, feature_names, feature_types, nthread)
    382         data, feature_names, feature_types = _maybe_pandas_data(data,
    383                                                                    feature
_names,
--> 384                                                                    feature
_types)
    385
    386         data, feature_names, feature_types = _maybe_dt_data(data,

~\Anaconda3\lib\site-packages\xgboost\core.py in _maybe_pandas_data(data, feature_names, feature_types)
    239         msg = ""
    240         Did not expect the data types in fields ""
--> 241         raise ValueError(msg + ', '.join(bad_fields))
    242
    243         if feature_names is None:

ValueError: DataFrame.dtypes for data must be int, float or bool.
Did not expect the data types in fields Name, Platform, Genre,
Publisher
```

```
In [80]: corr=Ensemble.corr()  
corr = (corr)  
sns.heatmap(corr,  
             xticklabels=corr.columns.values,  
             yticklabels=corr.columns.values)  
sns.plt.title('Heatmap of Correlation Matrix')
```

```
-----  
NameError                                Traceback (most recent call last)  
<ipython-input-80-35c59d31881a> in <module>  
----> 1 corr=Ensemble.corr()  
      2 corr = (corr)  
      3 sns.heatmap(corr,  
      4             xticklabels=corr.columns.values,  
      5             yticklabels=corr.columns.values)  
  
NameError: name 'Ensemble' is not defined
```

In []:

In []:

In []: